

Clustering

MGTF 495

Class Outline

- Representation Learning
 - k-means
 - EM optional
 - Agglomerative hierarchical clustering
 - Hands-On
- Informative Projections
 - PCA
 - SVD optional
 - Latent semantic indexing (LSI)
 - Hands-On

Representation learning



Good representations make learning easier.


Representation learning



Good representations make learning easier.

- They bring out the true degrees of freedom in the data.
- They capture relevant structure at multiple scales.
- They screen out noisy or irrelevant structure.


Degrees of freedom



Usual representation of speech:

- Take overlapping windows of the speech signal
- Apply many filters within each window
- More filters \Rightarrow higher dimensional

Degrees of freedom



Usual representation of speech:

- Take overlapping windows of the speech signal
- Apply many filters within each window
- More filters \Rightarrow higher dimensional

But the speech is produced by a physical system (vocal tract) with a fixed number of degrees of freedom. And the phoneme being uttered can be characterized by the configuration of this apparatus.

Multiscale structure



Commonly-occurring structure at many levels:

- Low-level: like local edges
- Higher-level: like wheels, windows

Representation learning: goals

How, and to what extent, can underlying degrees of freedom and multiscale structure be learned from the statistics of unlabeled data?

Representation learning: goals

How, and to what extent, can underlying degrees of freedom and multiscale structure be learned from the statistics of unlabeled data?

And when labels are available, how can a good representation be learned in tandem with the classifier?

Representation learning: goals


How, and to what extent, can underlying degrees of freedom and multiscale structure be learned from the statistics of unlabeled data?

And when labels are available, how can a good representation be learned in tandem with the classifier?


Topics:

- Clustering
- Informative linear projections
- Embedding and manifold learning
- Metric learning
- Autoencoders
- Deep nets


Clustering in \mathbf{R}^p



Clustering in \mathbf{R}^p




Clustering in \mathbf{R}^p



Two common uses of clustering:

Clustering in \mathbb{R}^p



Two common uses of clustering:

- **Vector quantization**

Find a finite set of representatives, that provides good coverage of a complex, possibly infinite, high-dimensional space

- **Finding meaningful structure in data**

Finding salient grouping in data.

Widely-used clustering methods

- ① K -means and its many variants
- ② EM for mixtures of Gaussians
- ③ Agglomerative hierarchical clustering


Class Outline

- Representation Learning
 - **k-means**
 - EM
 - Agglomerative hierarchical clustering
 - Hands-On
- Informative Projections
 - PCA
 - SVD
 - Latent semantic indexing (LSI)
 - Hands-On

The k -means optimization problem

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^p$; integer k
- Output: “Centers”, or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:


$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



The k -means optimization problem

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^p$; integer k
- Output: “Centers”, or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:


$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



The k -means optimization problem

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^p$; integer k
- Output: "Centers", or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:


$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



The k -means optimization problem

- Input: Points $x_1, \dots, x_n \in \mathbb{R}^p$; integer k
- Output: "Centers", or representatives, $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:

$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$




The centers carve \mathbb{R}^p up into k convex regions: μ_j 's region consists of points for which it is the closest center.

Lloyd's k -means algorithm

The k -means problem is NP-hard to solve. The most popular heuristic is called the “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

The k -means problem is NP-hard to solve. The most popular heuristic is called the “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

The k -means problem is NP-hard to solve. The most popular heuristic is called the “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

The k -means problem is NP-hard to solve. The most popular heuristic is called the “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

The k -means problem is NP-hard to solve. The most popular heuristic is called the “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

The k -means problem is NP-hard to solve. The most popular heuristic is called the “ k -means algorithm”.


- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Lloyd's k -means algorithm

The k -means problem is NP-hard to solve. The most popular heuristic is called the “ k -means algorithm”.

- Initialize centers μ_1, \dots, μ_k in some manner.
- Repeat until convergence:
 - Assign each point to its closest center.
 - Update each μ_j to the mean of the points assigned to it.



Each iteration reduces the cost \Rightarrow convergence to a local optimum.

Initializing the k -means algorithm

Typical practice: choose k data points at random as the initial centers.

Initializing the k -means algorithm

Typical practice: choose k data points at random as the initial centers.

Another common trick: start with extra centers, then prune later.

Initializing the k -means algorithm

Typical practice: choose k data points at random as the initial centers.

Another common trick: start with extra centers. then prune later.

A particularly good initializer: k -means++

- Pick a data point x at random as the first center
- Let $C = \{x\}$ (centers chosen so far)
- Repeat until desired number of centers is attained:
 - Pick a data point x at random from the following distribution:

$$\Pr(x) \propto \text{dist}(x, C)^2,$$

where $\text{dist}(x, C) = \min_{z \in C} \|x - z\|$

- Add x to C

Representing images using k -means codewords

Given a collection of images, how to represent as fixed-length vectors?



Representing images using k -means codewords

Given a collection of images, how to represent as fixed-length vectors?



patch of
fixed size

- Look at all $\ell \times \ell$ patches in all images. Extract features for each.
- Run k -means on this entire collection to get k centers.
- Now associate any image patch with its nearest center.
- Represent an image by a histogram over $\{1, 2, \dots, k\}$.

Representing images using k -means codewords

Given a collection of images, how to represent as fixed-length vectors?



patch of
fixed size

- Look at all $\ell \times \ell$ patches in all images. Extract features for each.
- Run k -means on this entire collection to get k centers.
- Now associate any image patch with its nearest center.
- Represent an image by a histogram over $\{1, 2, \dots, k\}$.

Such data sets are truly enormous.

Streaming and online computation

Streaming computation: for data sets that are too large to fit in memory.

- Make one pass (or maybe a few passes) through the data.
- On each pass:
 - See data points one at a time, in order.
 - Update models/parameters along the way.
- There is only enough space to store a tiny fraction of the data, or a perhaps short summary.

Streaming and online computation

Streaming computation: for data sets that are too large to fit in memory.

- Make one pass (or maybe a few passes) through the data.
- On each pass:
 - See data points one at a time, in order.
 - Update models/parameters along the way.
- There is only enough space to store a tiny fraction of the data, or a perhaps short summary.

Online computation: an even more lightweight setup, for data that is continuously being collected.

- Initialize a model.
- Repeat forever:
 - See a new data point.
 - Update model if need be.

Example: sequential k -means

- ① Set the centers μ_1, \dots, μ_k to the first k data points
- ② Set their counts to $n_1 = n_2 = \dots = n_k = 1$
- ③ Repeat, possibly forever:
 - Get next data point x
 - Let μ_j be the center closest to x
 - Update μ_j and n_j :

$$\mu_j = \frac{n_j \mu_j + x}{n_j + 1} \quad \text{and} \quad n_j = n_j + 1$$

K -means: the good and the bad

The good:

- Fast and easy.
- Effective in quantization.

The bad:

- Geared towards data in which the clusters are spherical, and of roughly the same radius.

K -means: the good and the bad

The good:

- Fast and easy.
- Effective in quantization.

The bad:

- Geared towards data in which the clusters are spherical, and of roughly the same radius.


Is there is a similarly-simple algorithm in which clusters of more general shape are accommodated?

Class Outline

- Representation Learning
 - k-means
 - EM
 - Agglomerative hierarchical clustering
 - Hands-On
- Informative Projections
 - PCA
 - SVD
 - Latent semantic indexing (LSI)
 - Hands-On


Hierarchical clustering

Choosing the number of clusters (k) is difficult.



Hierarchical clustering


Choosing the number of clusters (k) is difficult.



Often there is no single right answer, because of multiscale structure.

Hierarchical clustering


Choosing the number of clusters (k) is difficult.




Often there is no single right answer, because of multiscale structure.

Hierarchical clustering avoids these problems.

Example: gene expression data




The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm




- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm



- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

The single linkage algorithm



- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two clusters with the closest pair of points
- Disregard singleton clusters


Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two “closest” clusters

Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two “closest” clusters


How to measure the distance between two clusters of points, C and C' ?



Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
 - Merge the two “closest” clusters

How to measure the distance between two clusters of points, C and C' ?



- Single linkage

$$\text{dist}(C, C') = \min_{x \in C, x' \in C'} \|x - x'\|$$

- Complete linkage

$$\text{dist}(C, C') = \max_{x \in C, x' \in C'} \|x - x'\|$$

Average linkage

Three commonly-used variants:

- 1 Average pairwise distance between points in the two clusters

$$\text{dist}(C, C') = \frac{1}{|C| \cdot |C'|} \sum_{x \in C} \sum_{x' \in C'} \|x - x'\|$$

- 2 Distance between cluster centers

$$\text{dist}(C, C') = \|\text{mean}(C) - \text{mean}(C')\|$$

- 3 Ward's method: the increase in k -means cost occasioned by merging the two clusters

$$\text{dist}(C, C') = \frac{|C| \cdot |C'|}{|C| + |C'|} \|\text{mean}(C) - \text{mean}(C')\|^2$$

(penalize merging of large clusters)